

# Fashionable Learning

- A deep study on how networks learn



# Problem Formulation



## Understand how neural networks learn

Implement several networks to get a better insight into how they work and identify features in the classification of images



## Visualize the training

Use visualization tools to see how the networks change during training



## Find cutting-edge NNs

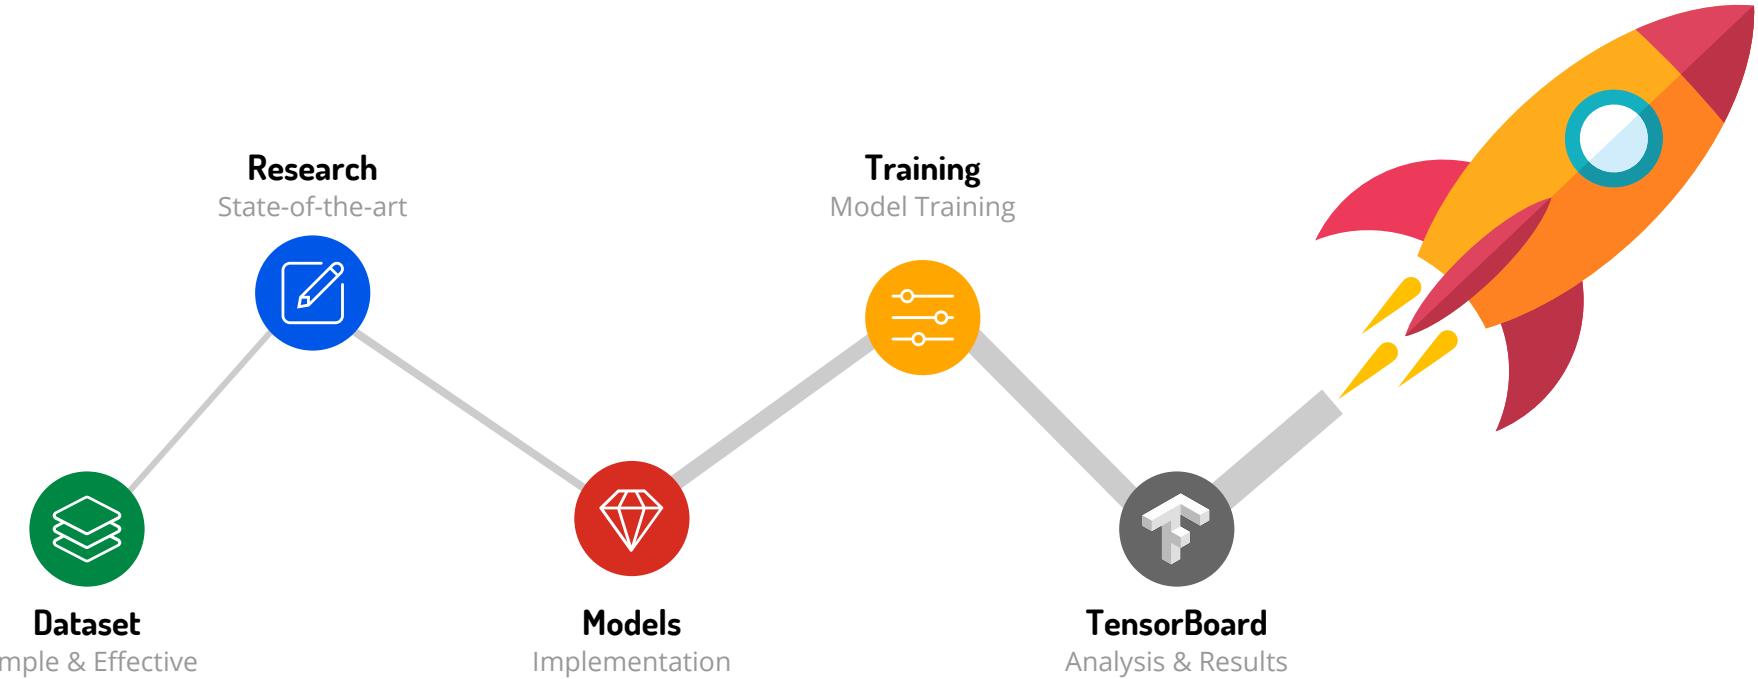
Compare our results with state-of-the art neural networks



## Dataset

Zalandoresearch's Fashion-MNIST - a more complicated MNIST-like dataset used for experimentation in deep learning

# Technical Approach





# Dataset

Fashion-MNIST

TRAINING

60 000

VALIDATION

5000

TEST

10 000

T-SNE



PCA



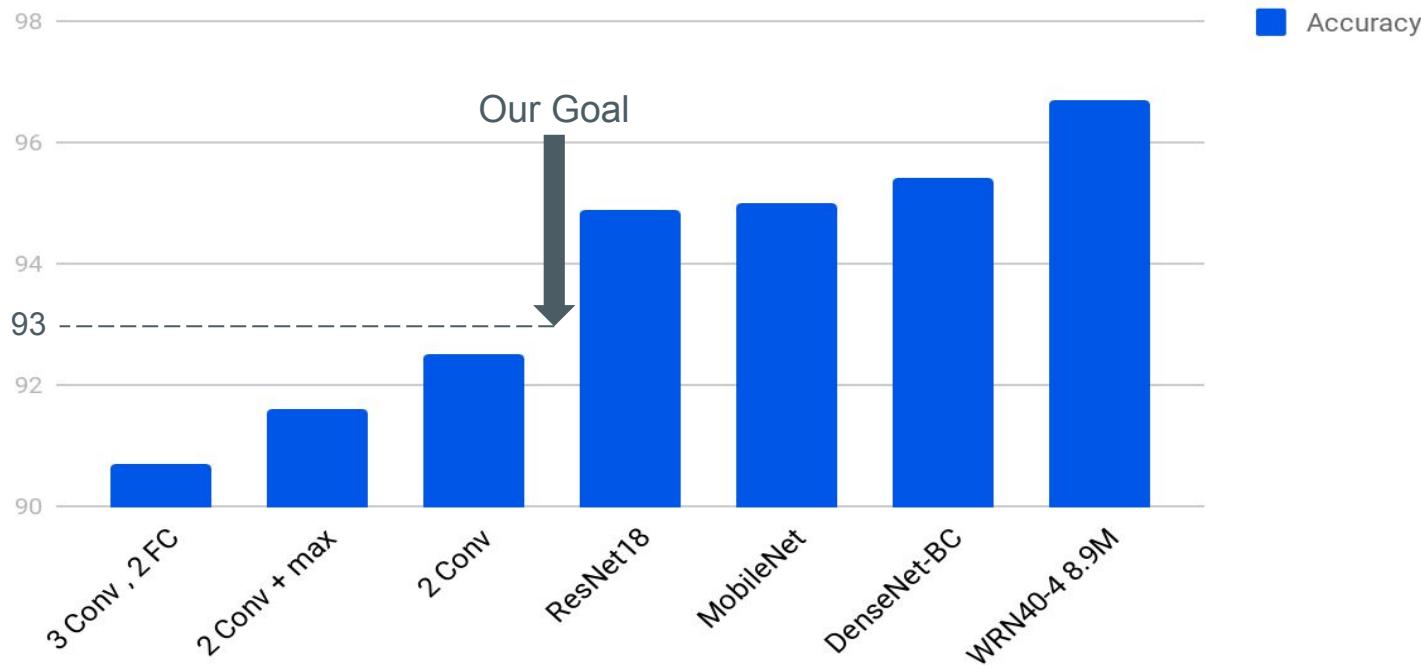
**Why did we choose this dataset?**  
**Simple** and **effective** for our purpose



# Research

A look into existing models, and a benchmarking of them

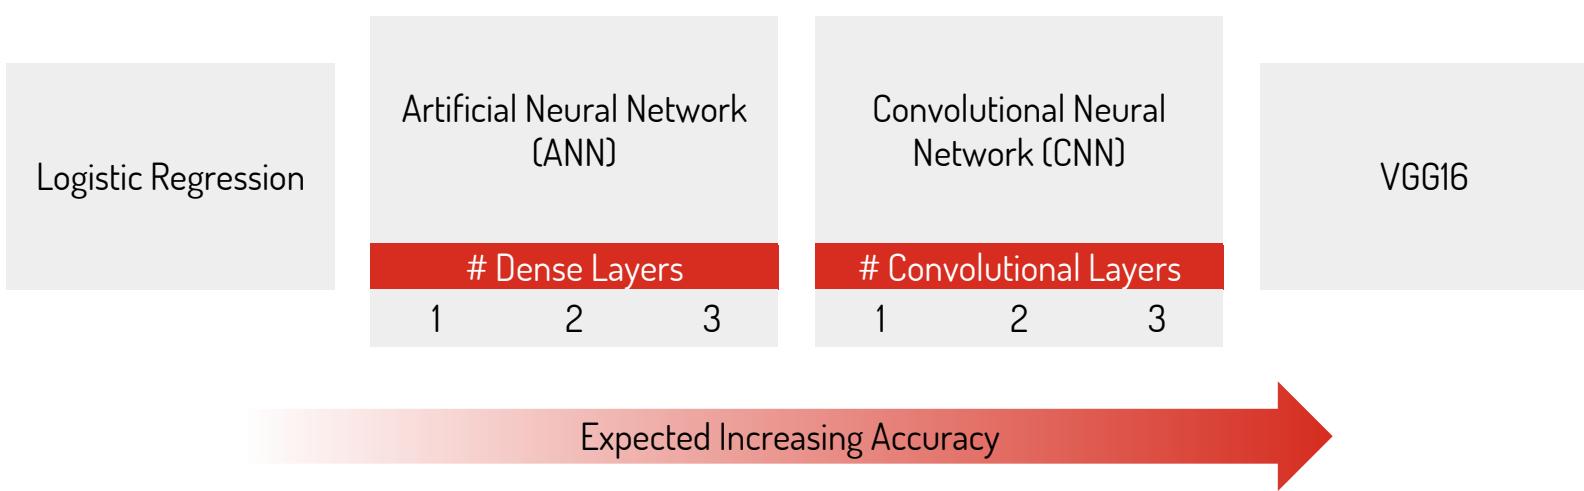
## Model performances





# Models

An overview of the models we chose to implement





# Training

## Tweaking and Tuning

EPOCHS

50

ITERATIONS

600

BATCH SIZE

100

Learning  
Rate

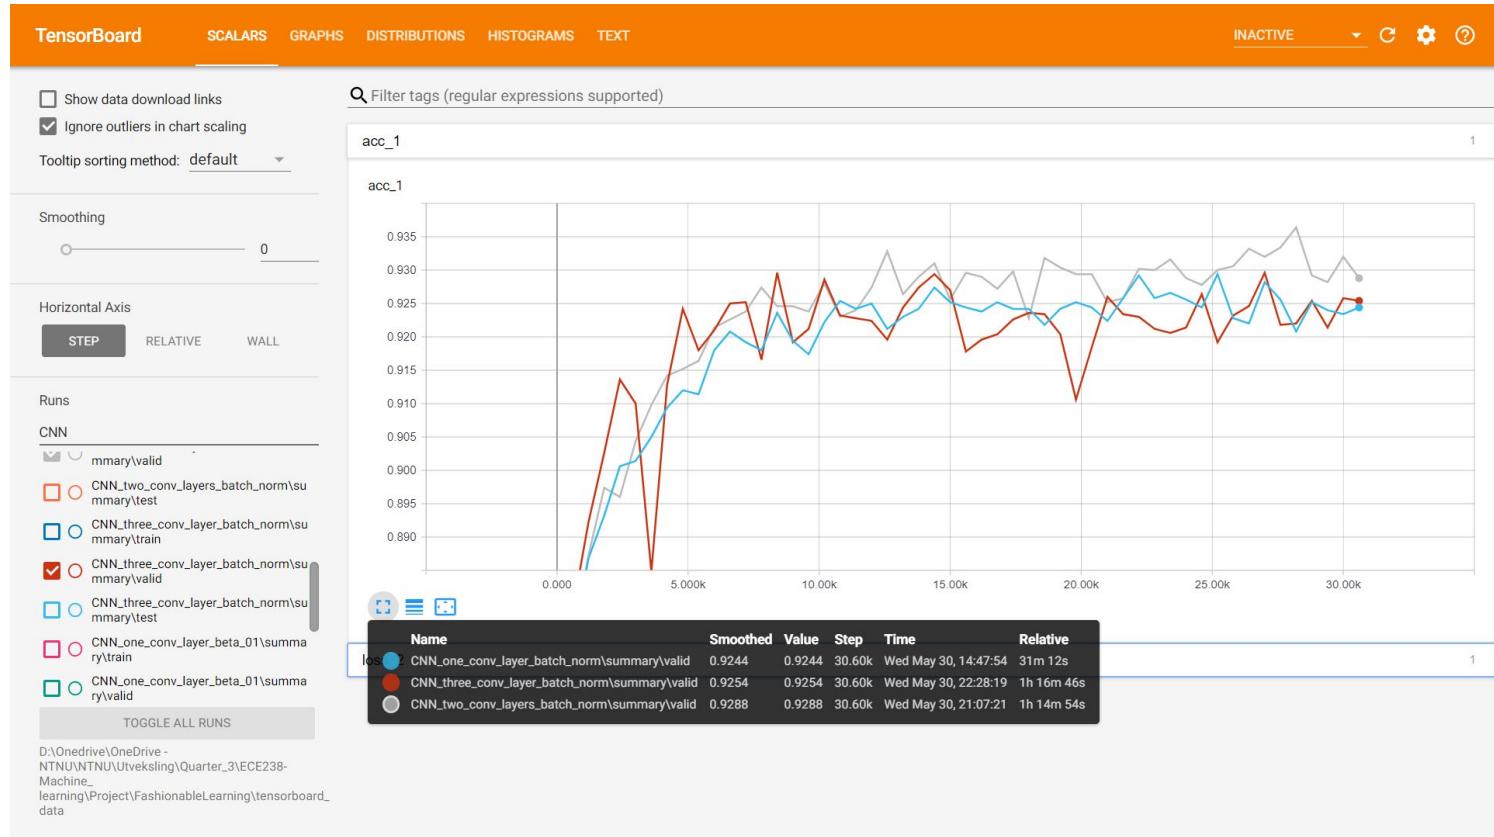
Regularization

Batch  
Normalization

Dropout  
Probability

# TensorBoard

TensorFlow's visualization tool





# ANALYSIS & RESULTS

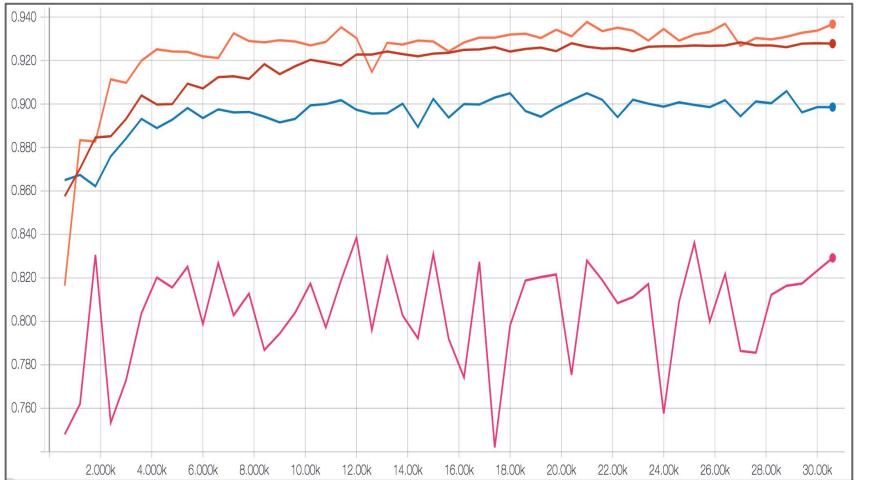
What did we find?

# TODO: REMOVE

## Analysis & Results Outline

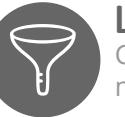
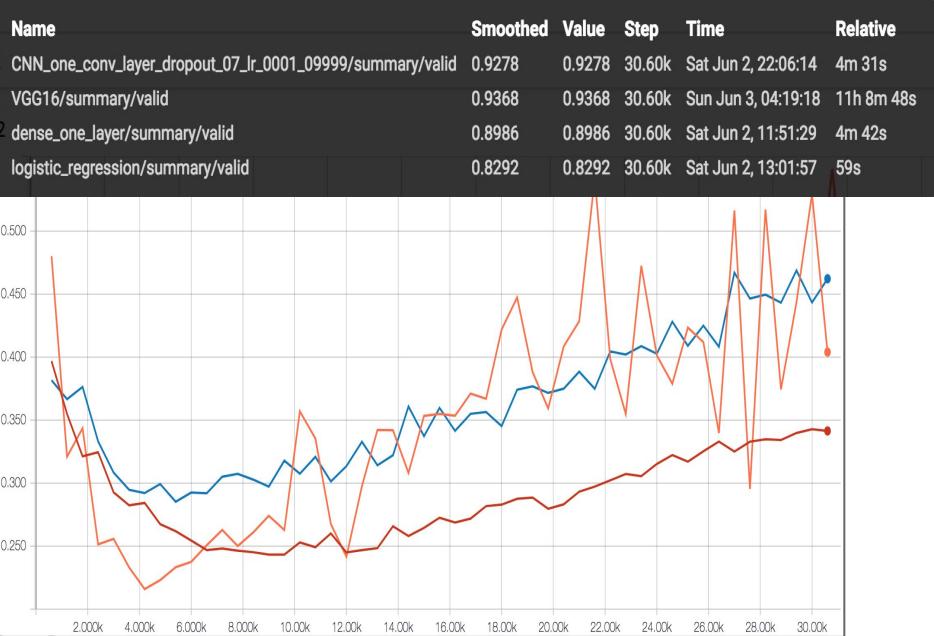
- Comparison
  - All networks (perhaps even with random forest as a representative for traditional machine learning?) (Acc + Loss)
  - The 3 ANNs (Acc + Loss)
  - The 3 CNNs
    - Acc + Loss
    - Histograms
    - Distributions
  - Accuracy vs computation time. Maybe compare the best ANN with the best CNN
- Deep-dive into the best-performing CNN (one layer)
  - Filter evolution (gif)
  - (Should we say something about how networks learn in general?)
  - How do the weights and biases (etc.) change over time?
  - Do all layers learn?
- Overfitting

# Comparison of all Networks



## Accuracy

VGG16 has the highest accuracy, as expected



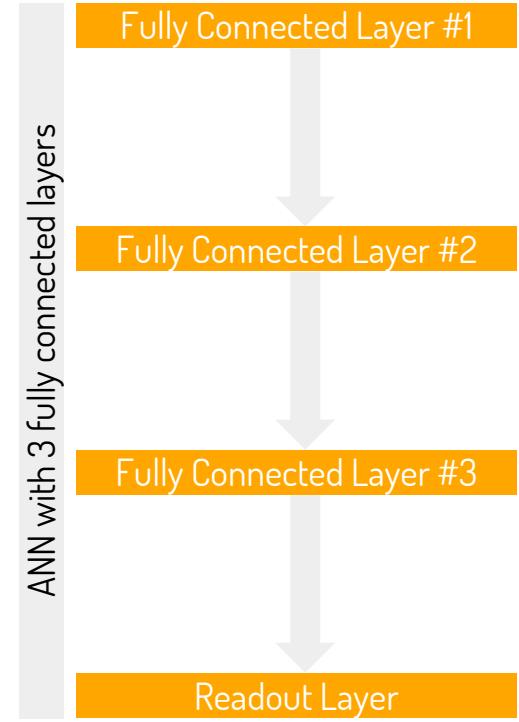
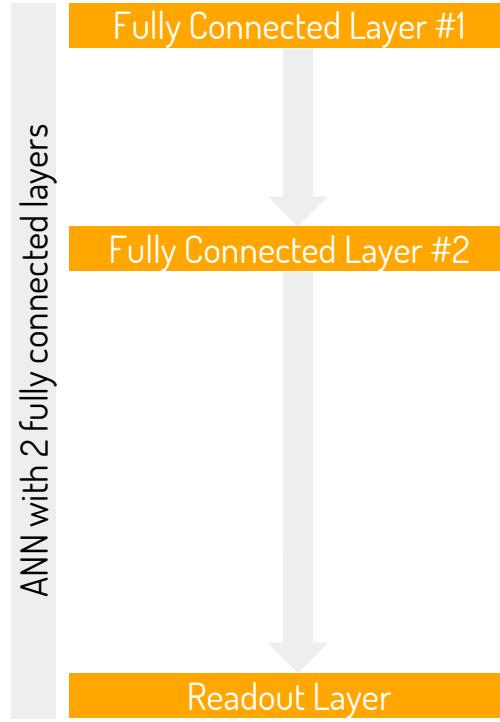
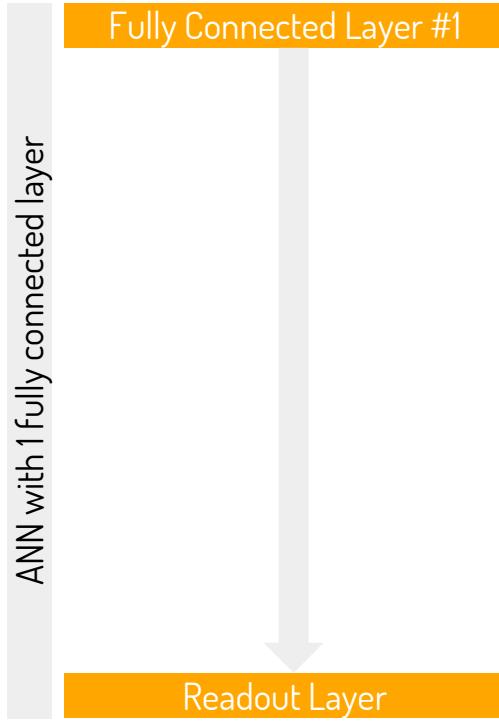
## Loss

CNN with one layer has the lowest loss, which is also more stable across epochs than the other networks.

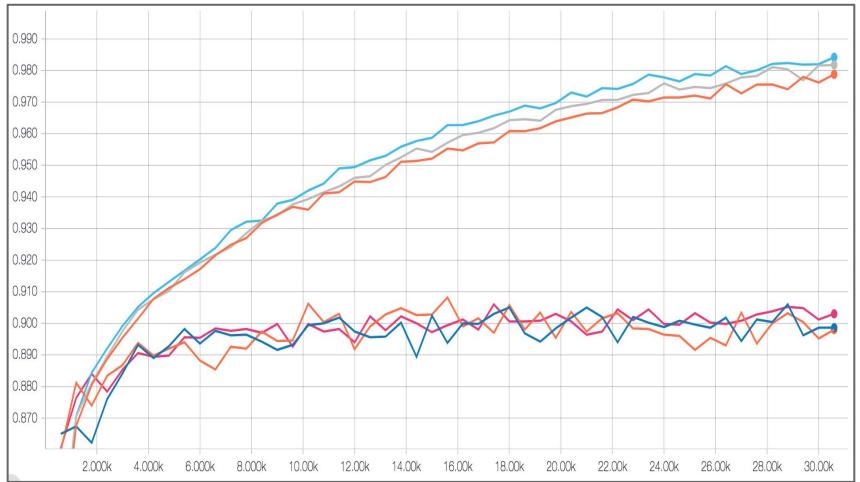
Test Accuracy			
LR	ANN	CNN	VGG16
0.8092	0.8960	0.9243	0.9265

Test Loss			
LR	ANN	CNN	VGG16
11.546	0.4950	0.3481	0.4840

# ANN comparisons (1/3)



# ANN comparisons (2/3)



## Accuracy

Negligible differences between each ANN



## Loss

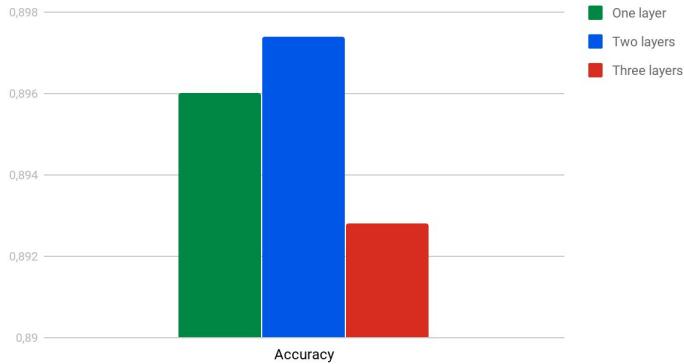
The ANNs with more layers have a higher loss  
Overfitting occurs after a few epochs

Test Accuracy		
1	2	3
0.8960	0.8974	0.8928

Test Loss		
1	2	3
0.4950	0.7346	0.7644

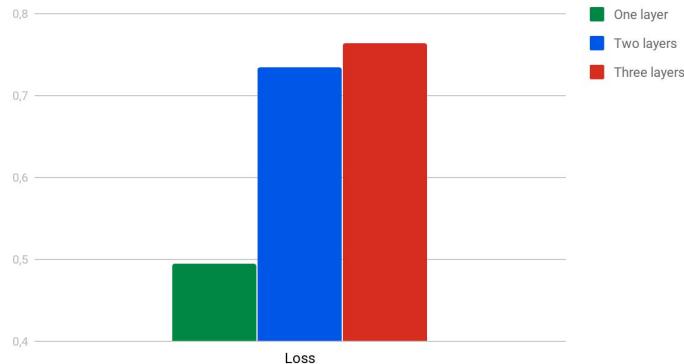
# ANN comparisons (3/3)

Test accuracy

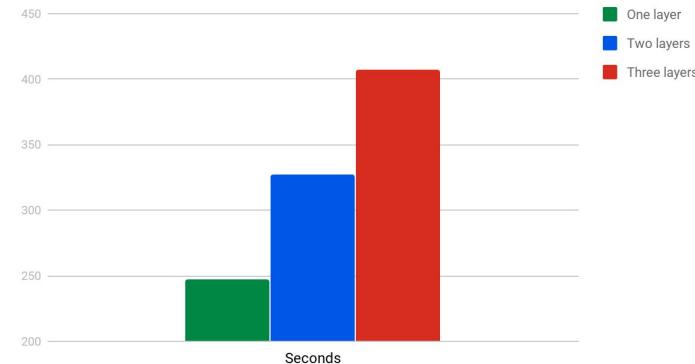


The one layer ANN is the neural network of choice (similar accuracy, and lower loss and comp. time than the others)

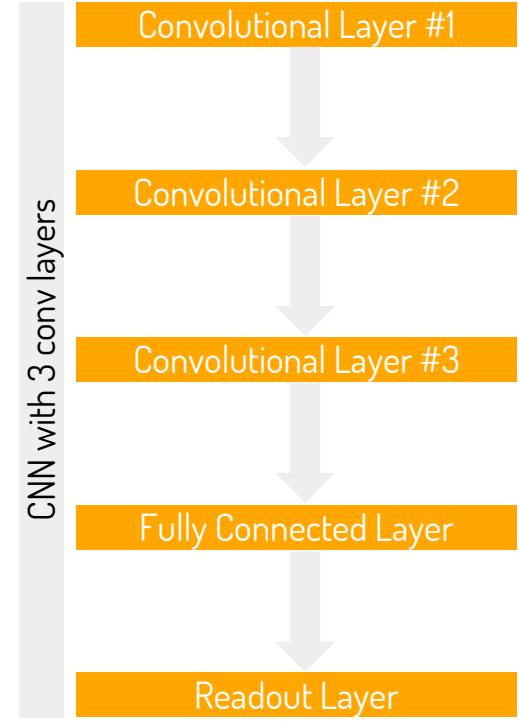
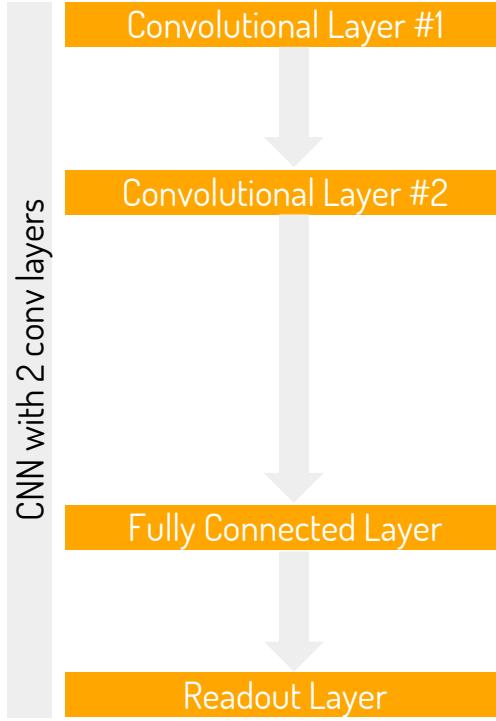
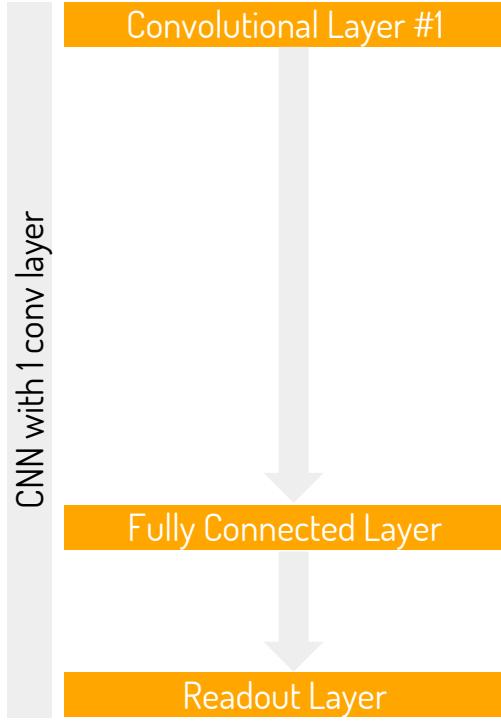
Test loss



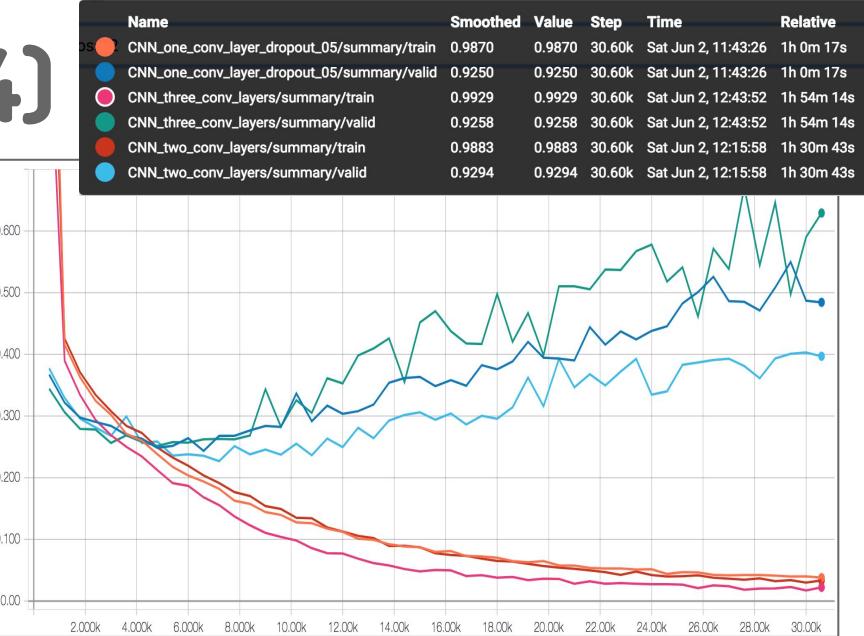
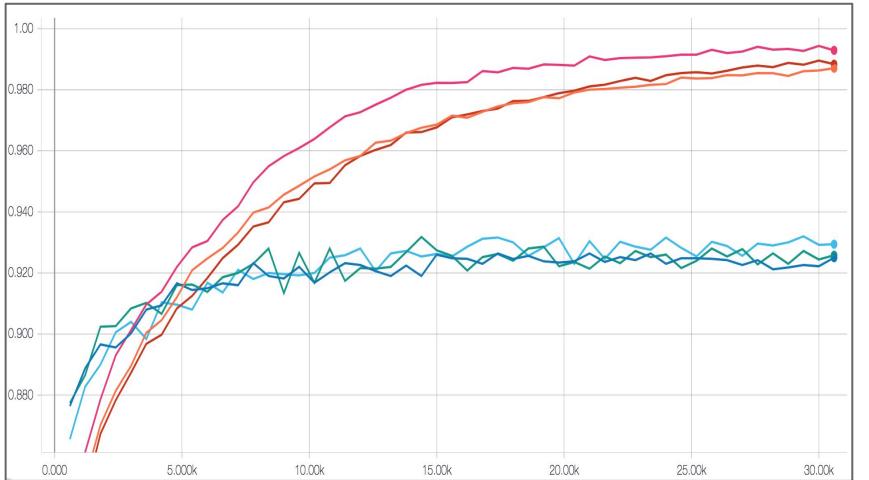
Computation time



# CNN comparisons (1/4)



# CNN comparisons (2/4)



## Accuracy

Negligible differences between each CNN



## Loss

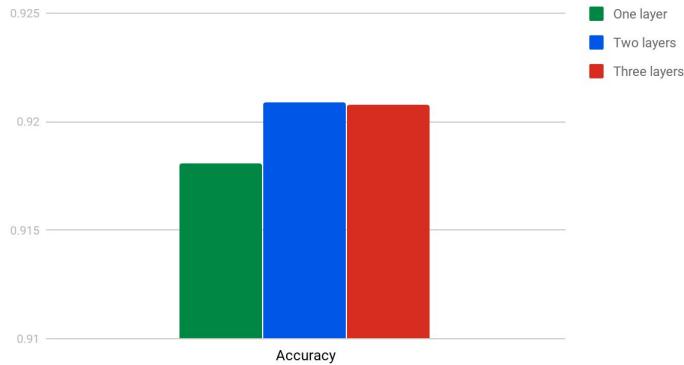
The CNN with two convolution layers has the smallest loss for validation data  
Overfitting occurs after a few epochs

Test Accuracy		
1	2	3
0.9181	0.9209	0.9208

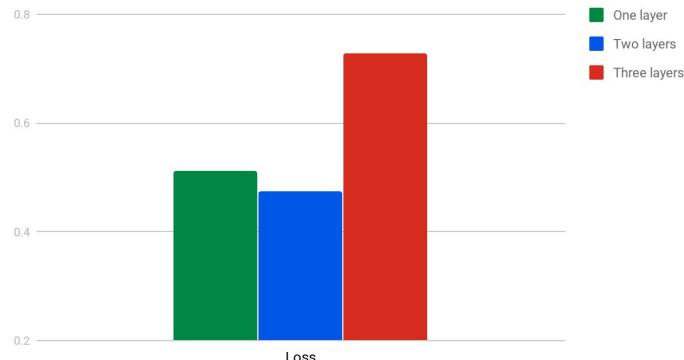
Test Loss		
1	2	3
0.5122	0.4740	0.7275

# CNN comparisons (3/4)

Test accuracy

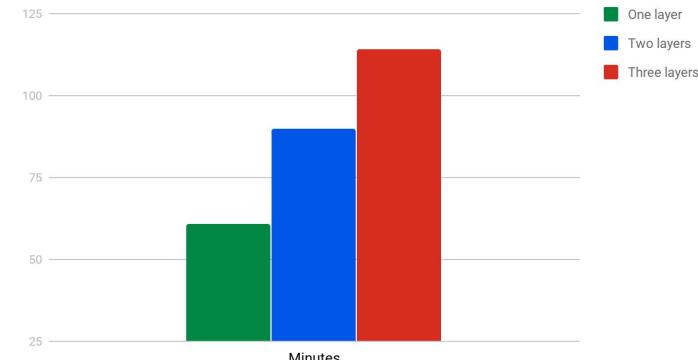


Test loss

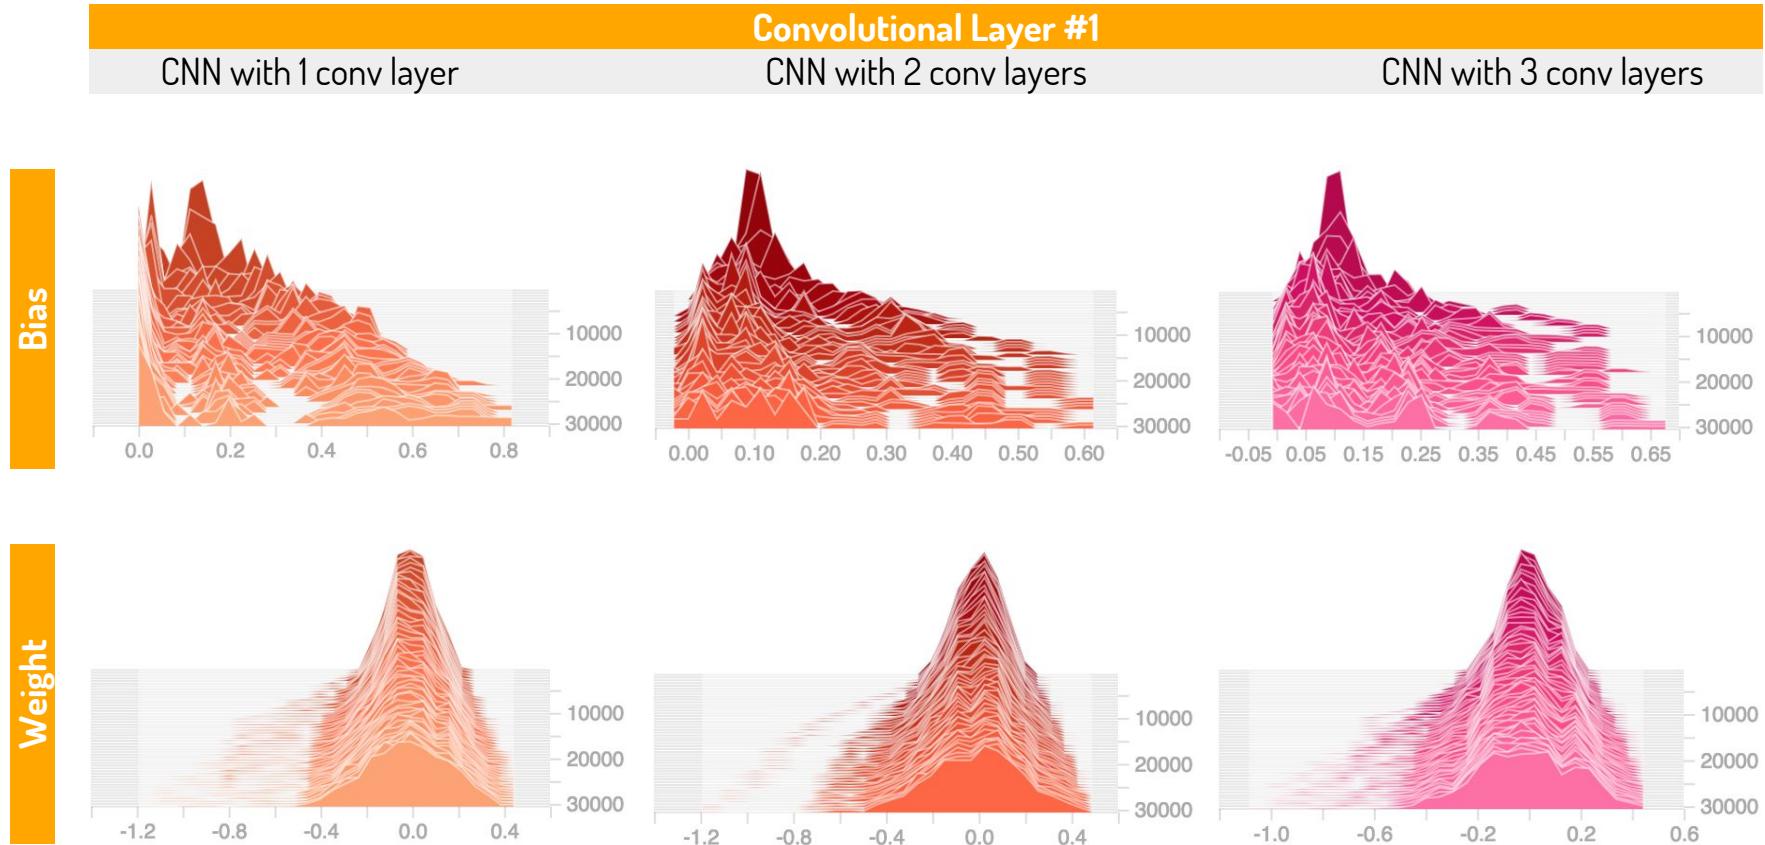


Even though the two-convolutional-layer CNN has higher accuracy and lower loss than the one-convolutional-layer, the difference is negligible compared to the difference in computation time. Hence, the one-convolutional-layer CNN is the CNN of choice.

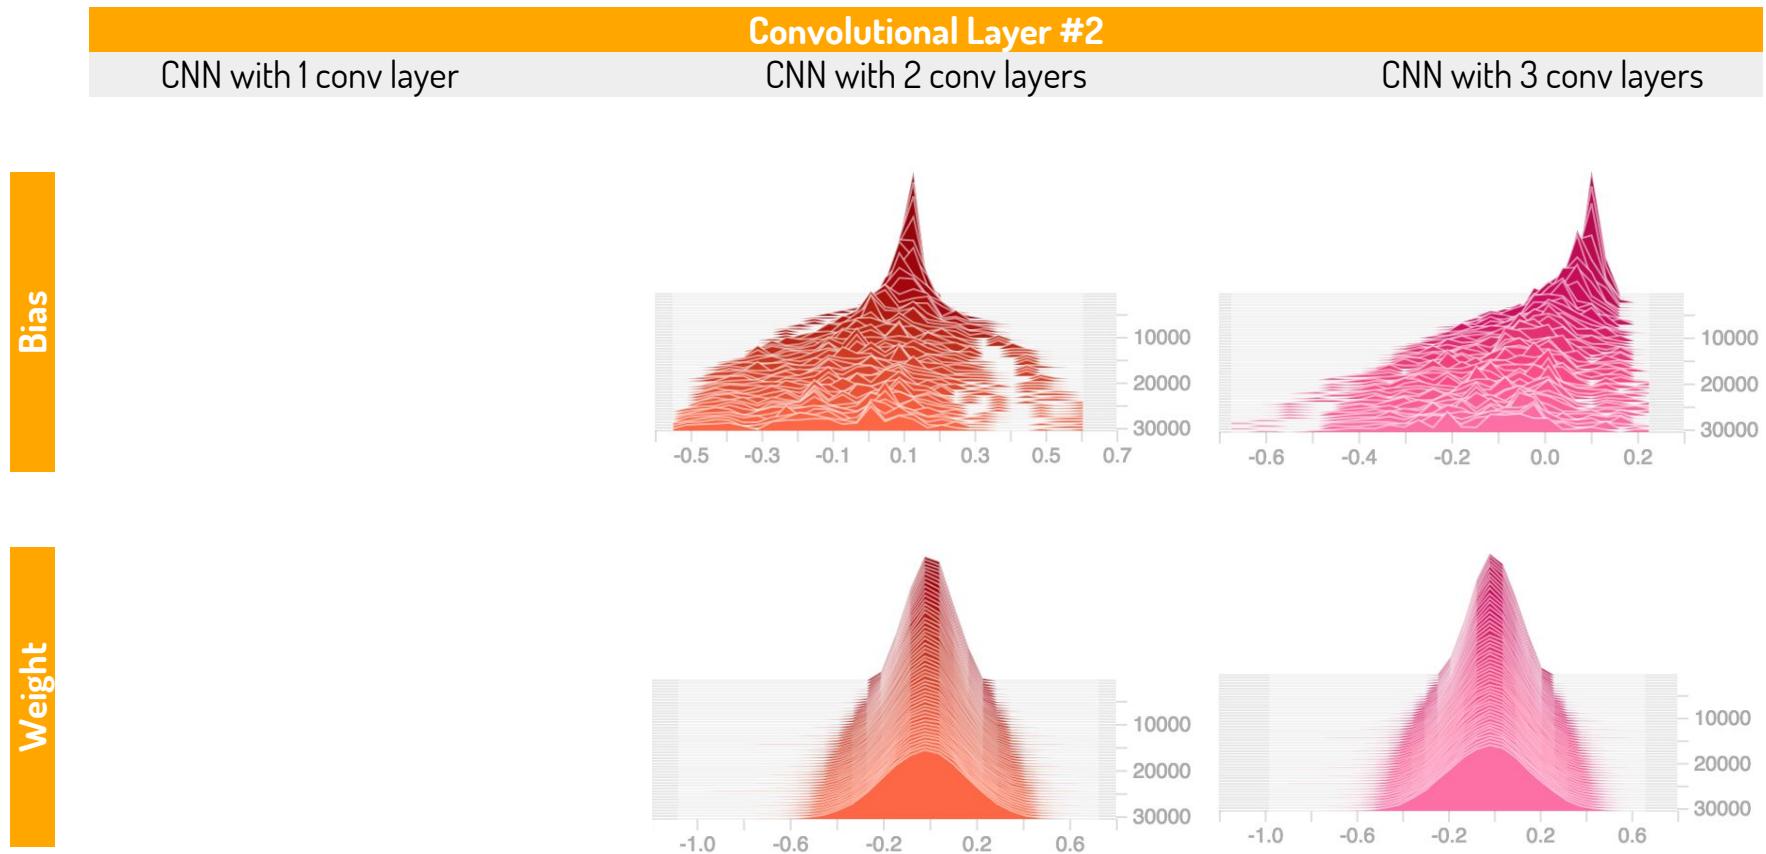
Computation time



# CNN comparisons (4/4) - Histograms



# CNN comparisons (4/4) - Histograms



# CNN comparisons (4/4) - Histograms

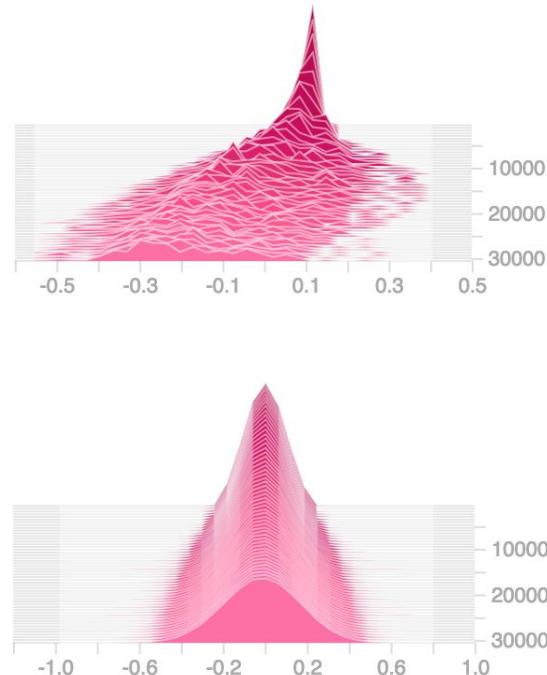
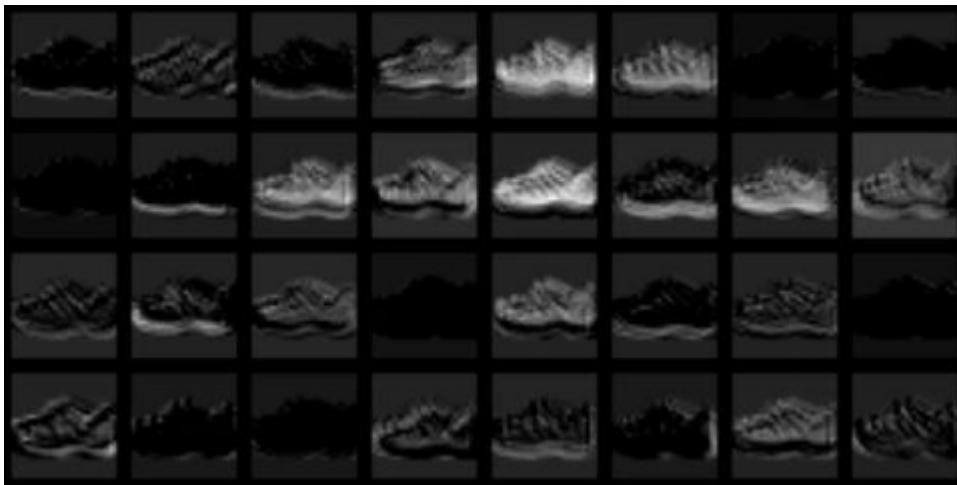
Bias  
Weight

Convolutional Layer #3

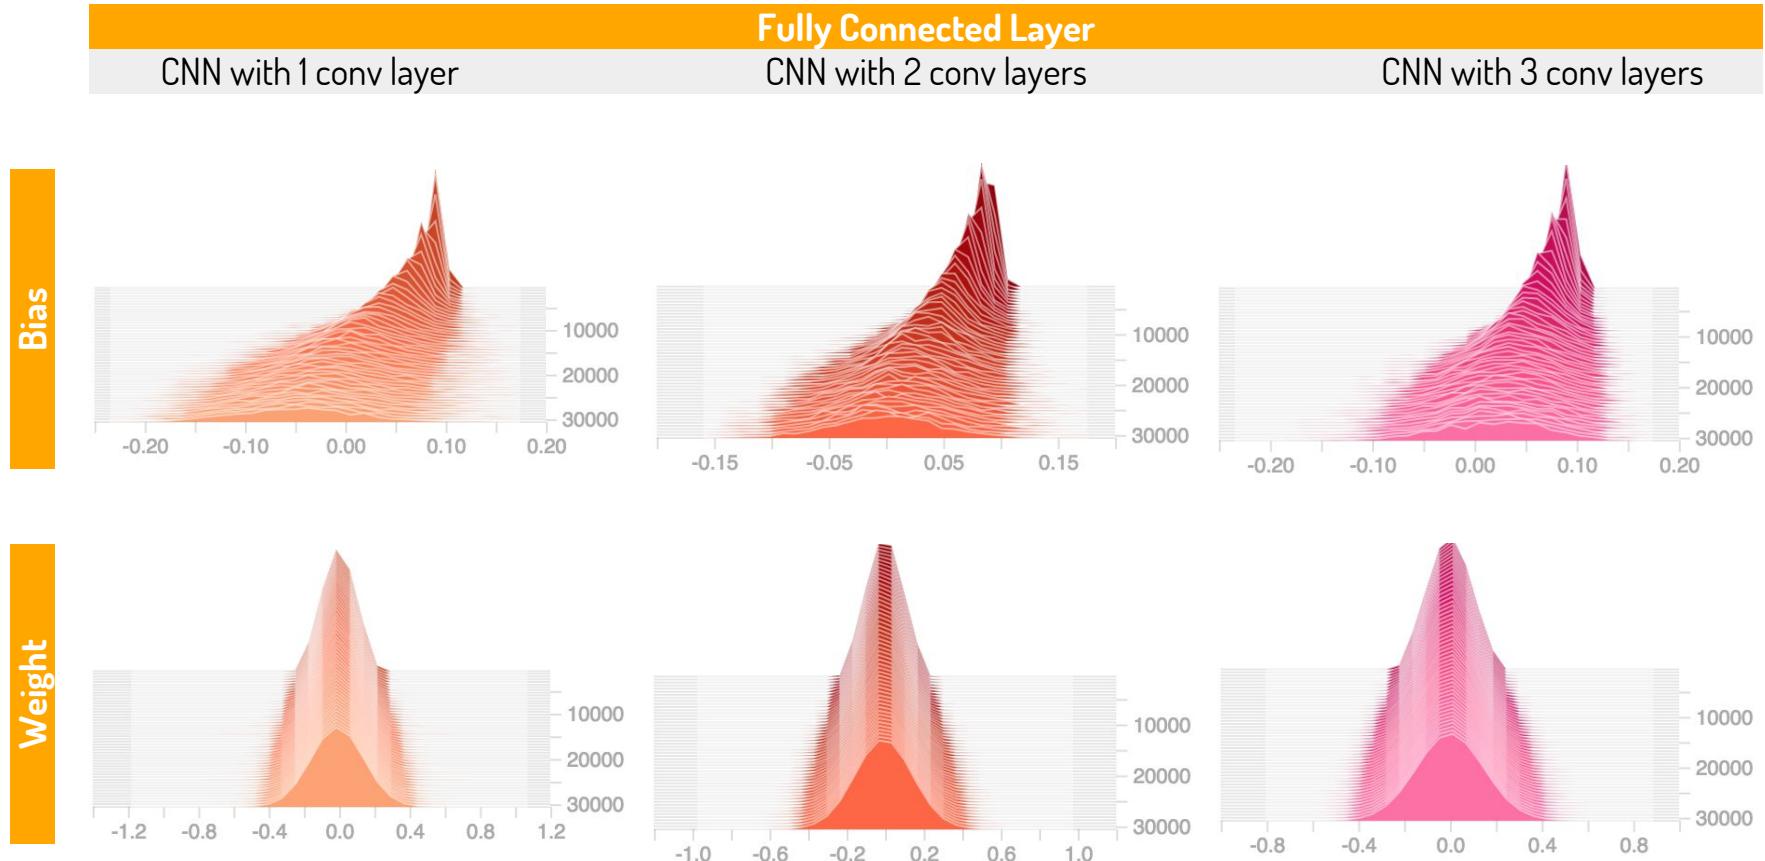
CNN with 1 conv layer

CNN with 2 conv layers

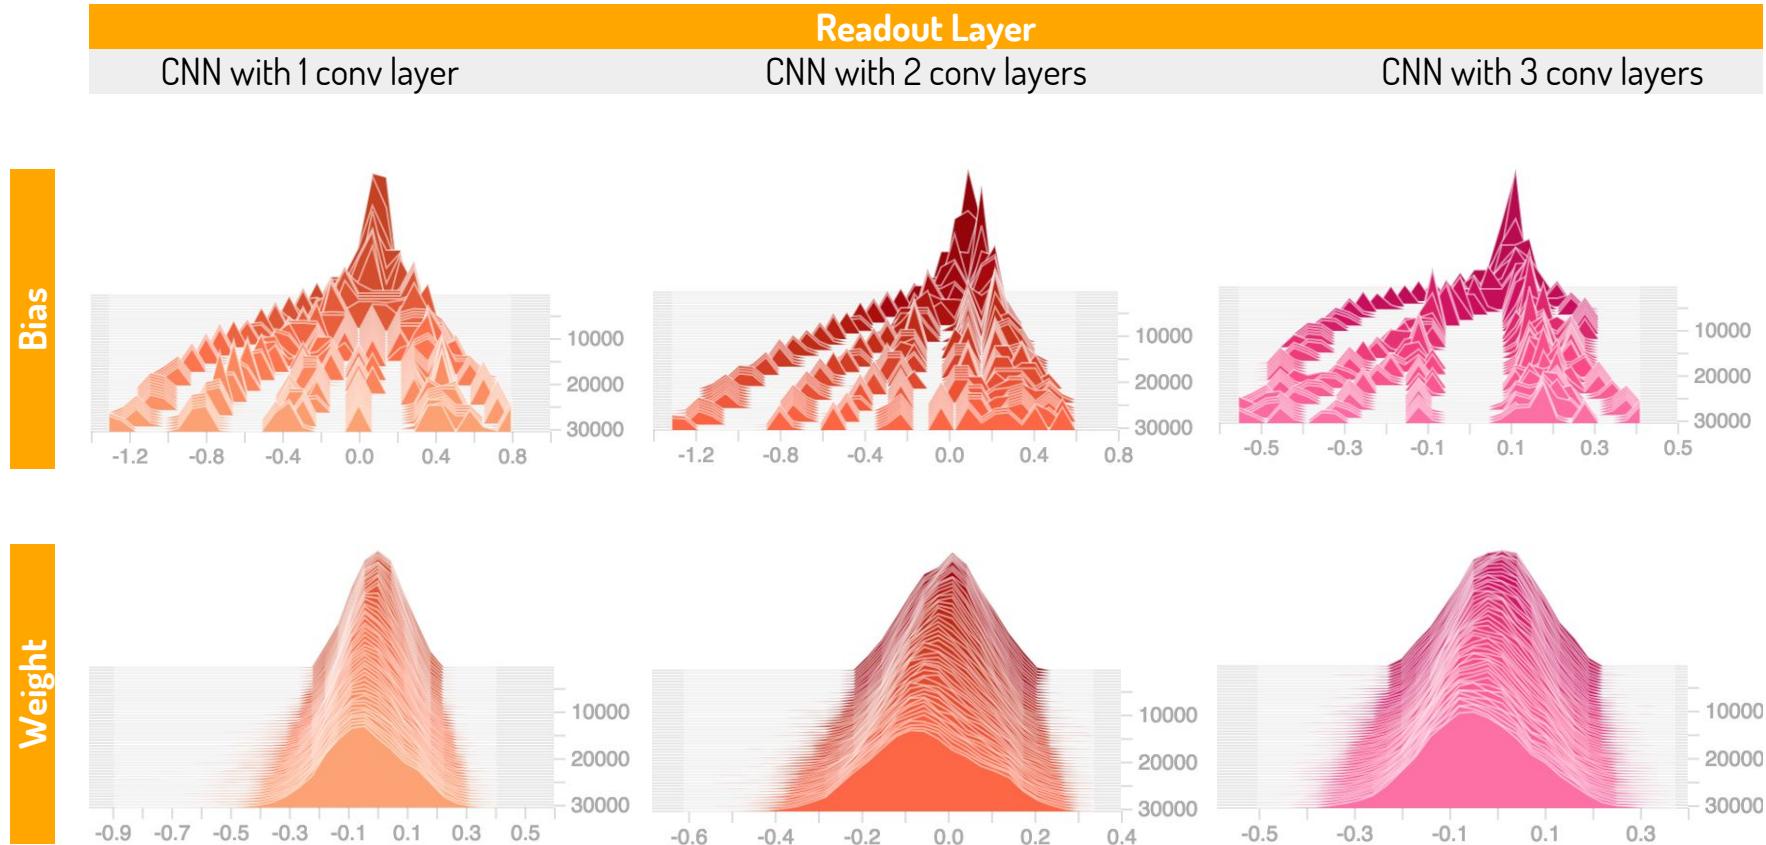
CNN with 3 conv layers



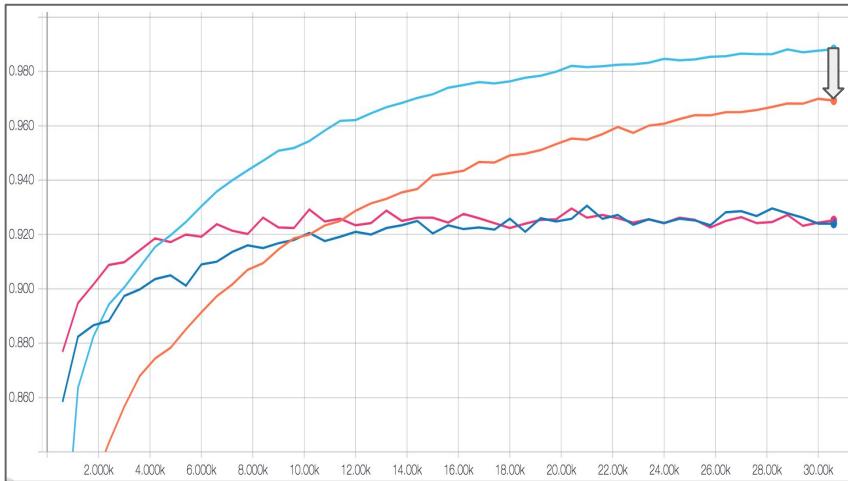
# CNN comparisons (4/4) - Histograms



# CNN comparisons (4/4) - Histograms

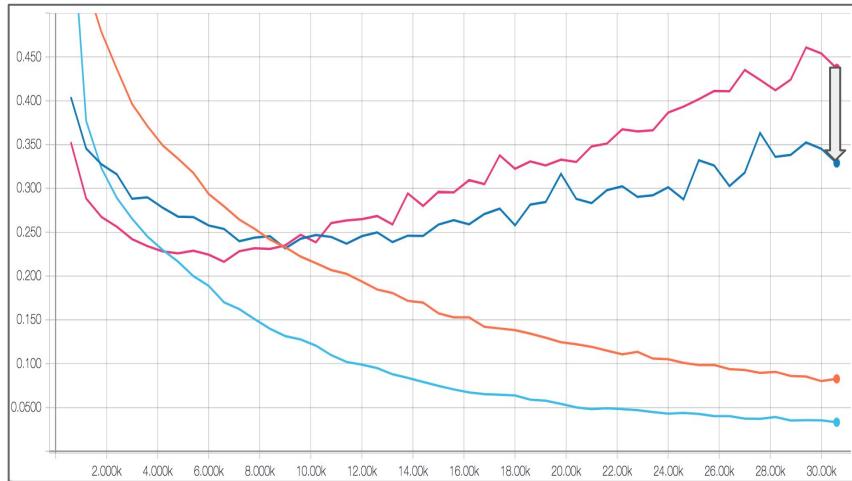


# Less Overfitting with Batch Normalization



## Accuracy

Very similar validation accuracy for both with and without batch normalization



## Loss

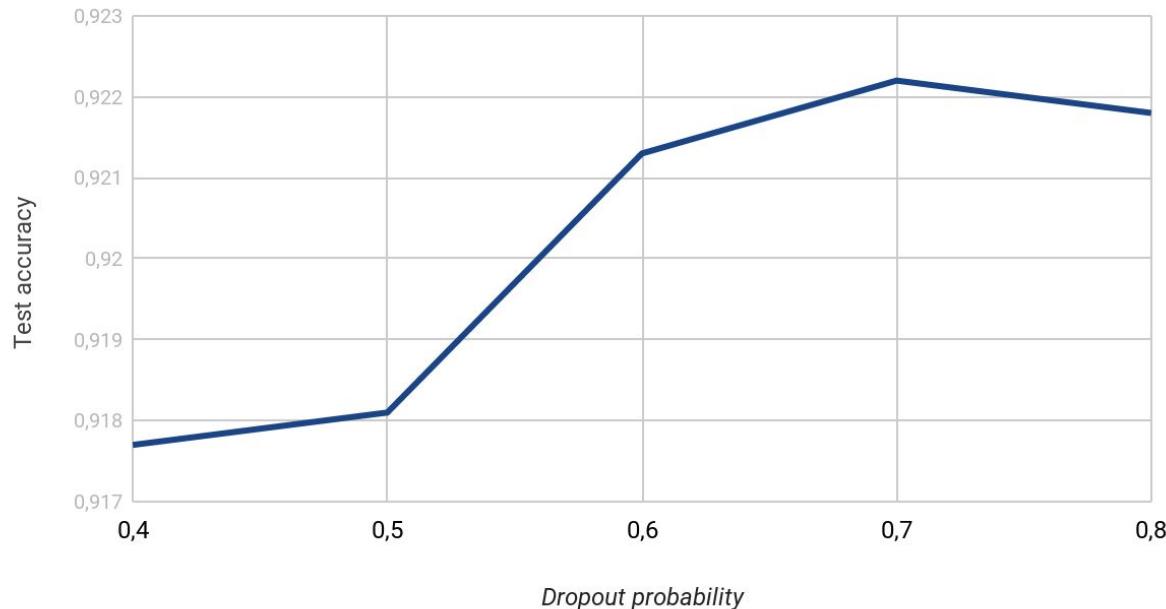
Batch normalization reduces the loss for the validation data, which implies less overfitting

Name	Smoothed	Value	Step	Time	Relative
CNN_one_conv_layer_dropout_07/summary/train	0.9693	0.9693	30.60k	Sat Jun 2, 19:43:52	4m 14s
CNN_one_conv_layer_dropout_07/summary/valid	0.9240	0.9240	30.60k	Sat Jun 2, 19:43:52	4m 14s
CNN_one_conv_layer_no_batch_norm/summary/train	0.9882	0.9882	30.60k	Sun Jun 3, 13:07:46	21m 34s
CNN_one_conv_layer_no_batch_norm/summary/valid	0.9252	0.9252	30.60k	Sun Jun 3, 13:07:46	21m 34s

# Dropout Probability

Can dramatically increase accuracy, with the added benefit of a lower loss

CNN with 1 conv layer



Dropout Probability

0,7

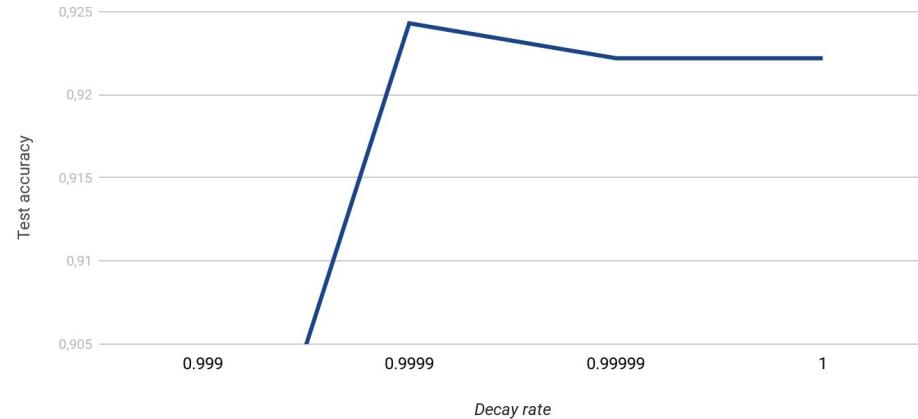
Max Test Accuracy

92,22%

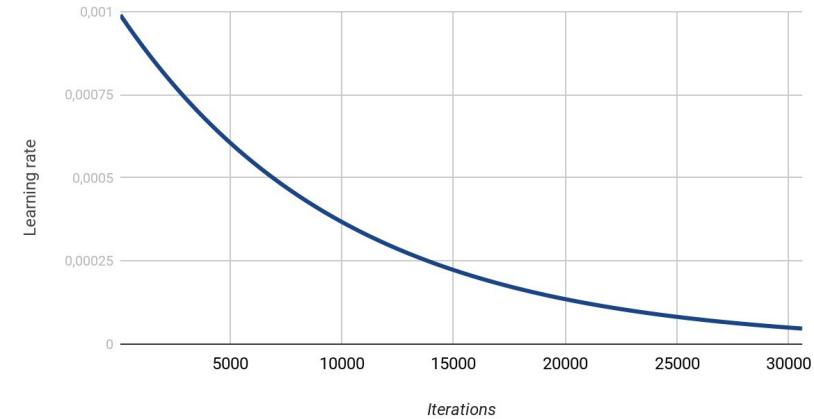
# Learning Rate

Lower steps might be needed to allow training, but at the risk of overfitting

CNN with 1 conv layer



Adaptive learning rate



Decay Rate

0.9999

Initial Learning Rate

0.001

Max Test Accuracy

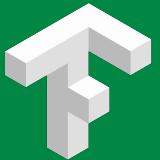
92.43%

$$L_r[k] = L_r[0] \cdot d_r^k$$



# LEARNING OUTCOME

Our takeaways and ideas



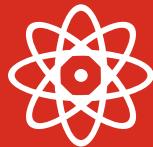
## TensorFlow + TensorBoard

Create neural networks of varying complexity



## Tuning and Tweaking

Learning rate, batch normalization, dropout probability, regularization, etc.



## Visualization

See how the networks evolve as they are being trained



## Analysis

Find the optimal model and training strategy for a given dataset



## A deeper insight

For future projects we will more easily see the needed type and complexity of a suitable network



THANK  

---

  
YOU